

ExAD

August 17, 2020

1 ExAD: An Ensemble Approach for Explanation-based Adversarial Detection

1.0.1 In this Jupyter notebook, we provide reference code for our whitebox evaluation.

For more details, please refer to Section 4.6 of the paper as well as Appendix E.

```
[1]: import os
# os.environ["CUDA_DEVICE_ORDER"]="PCI_BUS_ID"
# os.environ["CUDA_VISIBLE_DEVICES"]="1"

import tensorflow as tf
import keras
from keras import backend as K
from keras.models import Model
from keras.utils import to_categorical
from keras.layers import Input, Dense, Lambda, merge, Dropout, Flatten, Conv2D,
    MaxPooling2D
from keras.models import Sequential, model_from_json, load_model
from keras.layers.core import Activation
from keras import datasets
import matplotlib.image as mpimg
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
from numpy import clip
import pickle
import pandas as pd
from IPython.display import display
```

```
/Users/administrator/opt/anaconda3/envs/tf_env/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype(["qint8", np.int8, 1])
```

```
/Users/administrator/opt/anaconda3/envs/tf_env/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:524: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
```

```

numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
/Users/administrator/opt/anaconda3/envs/tf_env/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:525: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
/Users/administrator/opt/anaconda3/envs/tf_env/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:526: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
/Users/administrator/opt/anaconda3/envs/tf_env/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:527: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
/Users/administrator/opt/anaconda3/envs/tf_env/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:532: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_resource = np.dtype(["resource", np.ubyte, 1])
Using TensorFlow backend.

```

2 Utility Methods

```

[2]: # Utility Methods
def Union(lst1, lst2):
    final_list = list(set(lst1) | set(lst2))
    return final_list

```

3 Configuration

```

[3]: # Datasets on which we conduct our evaluation
datasets = ['mnist', 'fmnist', 'cifar10']

# The explanation methods that will be used in the ensemble
exp_methods = ['lrp', 'guided_backprop',
               'integrated_grad', 'pattern_attribution', 'grad_times_input']

# The directory of the detector models
# Format: 'data/defender/<dataset>/orig/train/<explanation_method>/
        '→<target_class_index>/model/'
model_dir_structure = 'data/defender/{}/orig/train/{}/{}/model/'
model_name = 'exp_model.h5'

```

```

# The directory where we store the adversarial examples created using whitebox_
→attack
# Format 'data/adv2/adversary/<dataset>/<target_exp_method>/
→<attack_method_to_test>/target_next/target_<class_idx>'
adv2_dir_structure = 'data/adversary/{}/adv2/{}/{} /target_next/target_{}/'
succ_examples_filename = 'succ_on_f.npy'

# The directory where we store the abnormal explanations created from_
→adversarial examples which were inturn created using whitebox attack
adv2_exp_dir_structure = 'data/adversary/{}/adv2/{}/{} /{} /from_{}/'
explanations_filename = 'expls.npy'

src_exp_methods = exp_methods[:]
src_exp_methods.append('overall')

enable_detailed_logs = False

```

4 Load the ExAD-CNN detector models

As discussed in Section 3.4.1 of our paper, for every class in a dataset, we train 5 detector models- one corresponding to every explanation technique.

```

[4]: import collections

# Dictionary of models.
# This avoids having to reload models multiple times, and therefore speeds up_
→the analysis.
print('\nLoading detector models')
if not enable_detailed_logs:
    print('Note: You have turned off option to view detailed logs. Loading_
→models may take a while.')

model_d = {}
for dataset in datasets:
    for class_idx in range(10):
        for exp_method in exp_methods:
            if enable_detailed_logs: print('\tLoading detector model for dataset:
→{} target class:{} explanation technique:{}'.format(dataset, class_idx, exp_method))
            model_dir = model_dir_structure.format(dataset, exp_method, _
→str(class_idx))
            model = load_model(model_dir + model_name)
            model_d[(dataset, class_idx, exp_method)] = model

```

Loading detector models

Note: You have turned off option to view detailed logs. Loading models may take a while.

```
[5]: df_list = list()

for dataset in datasets:
    print('Running whitebox attack evaluation on {} dataset'.format(dataset))

    # set the image dimension (side=H=W) based on the dataset
    if dataset in ['mnist', 'fmnist']:
        side = 28
    elif dataset in ['cifar10']:
        side = 32

    d = collections.defaultdict(list)
    total_adv_per_class_per_attack = 10

    table = []
    for target_exp_method in exp_methods:
        print('\n\tTargeting {} explanation technique'.format(target_exp_method))
        row = []
        for class_idx in range(10):
            if class_idx == 0:
                src_class_index = 9
            else:
                src_class_index = class_idx-1

            if target_exp_method == 'integrated_grad':
                attack_methods_to_test = ['cwl2/conf_0']
            else:
                attack_methods_to_test = ['cwl2/conf_0', 'cwlinf/conf_0', 'cwl0/
→conf_0', 'bim', 'mim', 'jsma']

            for attack_method_to_test in attack_methods_to_test:
                # The directory of adv2 examples which fool both f(.) and the
→target_exp_method
                # we do this to ensure we compute the performance on successful
→adv2 examples only
                adv2_dir = adv2_dir_structure.format(dataset, target_exp_method,
→attack_method_to_test, str(class_idx))
                succ_on_f = np.load(adv2_dir + succ_examples_filename)
                retained_adv_len = len(np.where(succ_on_f == True)[0])

                # if say 8/10 adversarial examples were successful, then failed
→will have indices from 0 to 7
```

```

failed = np.array([i for i in range(retained_adv_len)])

for exp_method in exp_methods:
    # Use the detector model for class_idx (target class)
    →corresponding to exp_method (explanation technique)
    model = model_d[(dataset, class_idx, exp_method)]

    adv2_exp_dir = adv2_exp_dir_structure.format(dataset,
    →target_exp_method, attack_method_to_test, exp_method, str(src_class_index) )
    adv2_exp = np.load(adv2_exp_dir + explanations_filename)

    # NOTE we retain successful adv2 examples
    adv2_exp = adv2_exp[succ_on_f]

    # process images for classification
    adv2_exp *= 255.0/np.max(adv2_exp)
    adv2_exp = adv2_exp.astype(int)
    adv2_exp = adv2_exp.reshape(-1,side,side,1)

    #evaluate model on adv2_exp samples
    result_test = model.predict(adv2_exp)

    result_test_class = np.argmax(result_test, axis=1)
    true_pos = len(np.where(result_test_class==1)[0])
    total_pos = len(result_test_class)
    det_rate = true_pos*100/total_pos

    failed_cur_method = np.where(result_test_class==0)[0]
    failed = np.intersect1d(failed, failed_cur_method)

    d[(target_exp_method, exp_method)].append(det_rate)
    d[(target_exp_method, attack_method_to_test, exp_method)].
    →append(det_rate)

    true_pos_cumulative = total_adv_per_class_per_attack -
    →len(failed)
    det_rate_cumulative = true_pos_cumulative * 100 /
    →total_adv_per_class_per_attack
    d[(target_exp_method, 'overall')].append(det_rate_cumulative)
    d[(target_exp_method, attack_method_to_test, 'overall')].
    →append(det_rate_cumulative)

for src_exp_method in src_exp_methods:
    l = np.array(d[(target_exp_method, src_exp_method)])
    mean_detection_rate = round(np.mean(l),3)

```

```

        row.append(mean_detection_rate)

    table.append(row)

arr = np.array(table)
df = pd.DataFrame(arr, index=exp_methods, columns=src_exp_methods)
df_list.append(df)

print('\n\nThis is the whitebox results for {}. \nThese results should be
→nearly consistent with Figure 8 of the paper in Appendix.'.format(dataset))
print('First column shows the targeted explanaton technique. Columns 1-5
→shows detection results by detector models corresponding to the explanation
→technique. Column 6 (rightmost) shows overall detection results by ExAD.')
print(df)

```

Running whitebox attack evaluation on mnist dataset

Targeting lrp explanation technique

Targeting guided_backprop explanation technique

Targeting integrated_grad explanation technique

Targeting pattern_attribution explanation technique

Targeting grad_times_input explanation technique

This is the whitebox results for mnist.

These results should be nearly consistent with Figure 8 of the paper in Appendix.

First column shows the targeted explanaton technique. Columns 1-5 shows detection results by detector models corresponding to the explanation technique. Column 6 (rightmost) shows overall detection results by ExAD.

	lrp	guided_backprop	integrated_grad	\
lrp	0.000	22.667	82.741	
guided_backprop	10.444	0.000	84.245	
integrated_grad	100.000	98.750	18.095	
pattern_attribution	1.542	31.476	81.077	
grad_times_input	98.722	99.458	56.700	

	pattern_attribution	grad_times_input	overall
lrp	15.833	83.463	90.667
guided_backprop	7.481	85.838	89.000
integrated_grad	90.000	55.837	100.000
pattern_attribution	10.167	85.276	89.667
grad_times_input	89.111	14.111	100.000

Running whitebox attack evaluation on fmnist dataset

Targeting lrp explanation technique

Targeting guided_backprop explanation technique

Targeting integrated_grad explanation technique

Targeting pattern_attribution explanation technique

Targeting grad_times_input explanation technique

This is the whitebox results for fmnist.

These results should be nearly consistent with Figure 8 of the paper in Appendix.

First column shows the targeted explanaton technique. Columns 1-5 shows detection results by detector models corresponding to the explanation technique. Column 6 (rightmost) shows overall detection results by ExAD.

	lrp	guided_backprop	integrated_grad	\
lrp	0.000	27.111	86.148	
guided_backprop	55.378	6.868	90.060	
integrated_grad	100.000	96.667	41.000	
pattern_attribution	13.720	34.176	84.880	
grad_times_input	95.394	95.033	71.955	

	pattern_attribution	grad_times_input	overall
lrp	20.000	93.130	95.667
guided_backprop	28.648	95.759	96.833
integrated_grad	96.667	89.000	100.000
pattern_attribution	12.870	93.403	95.500
grad_times_input	95.871	41.372	99.333

Running whitebox attack evaluation on cifar10 dataset

Targeting lrp explanation technique

Targeting guided_backprop explanation technique

Targeting integrated_grad explanation technique

Targeting pattern_attribution explanation technique

Targeting grad_times_input explanation technique

This is the whitebox results for cifar10.

These results should be nearly consistent with Figure 8 of the paper in Appendix.

First column shows the targeted explanaton technique. Columns 1-5 shows detection results by detector models corresponding to the explanation technique.

Column 6 (rightmost) shows overall detection results by ExAD.

	lrp	guided_backprop	integrated_grad	\
lrp	26.111	3.685	58.898	
guided_backprop	33.106	10.167	63.088	
integrated_grad	97.000	71.972	12.111	
pattern_attribution	25.773	7.940	60.394	
grad_times_input	79.398	62.116	50.264	

	pattern_attribution	grad_times_input	overall
lrp	50.000	57.921	84.000
guided_backprop	66.991	56.787	88.833
integrated_grad	90.778	49.111	99.000
pattern_attribution	24.444	50.023	80.667
grad_times_input	89.370	2.759	97.500

5 Review the results of whitebox attack

5.1 Review the results of whitebox attack on individual datasets

The results below should be consistent with Figure 8 of the paper in Appendix E.

How to interpret the results: First (or index) column shows the targeted explanaton technique. Columns 1-5 shows detection results by detector models corresponding to different explanation techniques. Column 6 (rightmost) shows overall detection results by ExAD (under CNN detector model setting).

```
[6]: for dataset_idx in range(len(df_list)):
      print('\n\nWhitebox results for {}'.format(datasets[dataset_idx]))
      display(df_list[dataset_idx])
```

Whitebox results for mnist.

	lrp	guided_backprop	integrated_grad	\
lrp	0.000	22.667	82.741	
guided_backprop	10.444	0.000	84.245	
integrated_grad	100.000	98.750	18.095	
pattern_attribution	1.542	31.476	81.077	
grad_times_input	98.722	99.458	56.700	

	pattern_attribution	grad_times_input	overall
lrp	15.833	83.463	90.667
guided_backprop	7.481	85.838	89.000
integrated_grad	90.000	55.837	100.000
pattern_attribution	10.167	85.276	89.667
grad_times_input	89.111	14.111	100.000

Whitebox results for fmnist.

	lrp	guided_backprop	integrated_grad	\
lrp	0.000	27.111	86.148	
guided_backprop	55.378	6.868	90.060	
integrated_grad	100.000	96.667	41.000	
pattern_attribution	13.720	34.176	84.880	
grad_times_input	95.394	95.033	71.955	
	pattern_attribution	grad_times_input	overall	
lrp	20.000	93.130	95.667	
guided_backprop	28.648	95.759	96.833	
integrated_grad	96.667	89.000	100.000	
pattern_attribution	12.870	93.403	95.500	
grad_times_input	95.871	41.372	99.333	

Whitebox results for cifar10.

	lrp	guided_backprop	integrated_grad	\
lrp	26.111	3.685	58.898	
guided_backprop	33.106	10.167	63.088	
integrated_grad	97.000	71.972	12.111	
pattern_attribution	25.773	7.940	60.394	
grad_times_input	79.398	62.116	50.264	
	pattern_attribution	grad_times_input	overall	
lrp	50.000	57.921	84.000	
guided_backprop	66.991	56.787	88.833	
integrated_grad	90.778	49.111	99.000	
pattern_attribution	24.444	50.023	80.667	
grad_times_input	89.370	2.759	97.500	

5.2 Review aggregated results of whitebox attack across datasets

The results below should be consistent with Figure 4 of the paper.

How to interpret the results: First (or index) column shows the targeted explanaton technique. Columns 1-5 shows detection results by detector models corresponding to different explanation techniques. Column 6 (rightmost) shows overall detection results by ExAD (under CNN detector model setting). Note that each cell reports a value computed by taking element-wise mean over each dataset. For e.g., in row 1 column 3, we show the mean detection rate (across datasets) corresponding to targeting LRP technique and using the detector model corresponding to GBP technique.

```
[7]: df1 = df_list[0]
      df2 = df_list[1]
      df3 = df_list[2]

      from functools import reduce

      dfs = [df1, df2, df3]
      df = reduce(lambda x, y: x.add(y), dfs) / len(dfs)
      display(df)
```

	lrp	guided_backprop	integrated_grad	\
lrp	8.703667	17.821000	75.929000	
guided_backprop	32.976000	5.678333	79.131000	
integrated_grad	99.000000	89.129667	23.735333	
pattern_attribution	13.678333	24.530667	75.450333	
grad_times_input	91.171333	85.535667	59.639667	

	pattern_attribution	grad_times_input	overall
lrp	28.611000	78.171333	90.111333
guided_backprop	34.373333	79.461333	91.555333
integrated_grad	92.481667	64.649333	99.666667
pattern_attribution	15.827000	76.234000	88.611333
grad_times_input	91.450667	19.414000	98.944333

```
[ ]:
```